

```
try:
    import google.colab
    IN_COLAB = True
except:
    IN_COLAB = False
if IN_COLAB:
    print("Downloading Colab files")
    ! shred -u setup_google_colab.py
    ! wget https://raw.githubusercontent.com/hse-aml/bayesian-methods-for-ml/master/setup_google_colab.py
    import setup_google_colab
    setup_google_colab.load_data_week5()
```

```
[>] Downloading Colab files
shred: setup_google_colab.py: failed to open for writing: No such file or directory
--2019-06-12 09:06:23-- https://raw.githubusercontent.com/hse-aml/bayesian-methods-for-ml/master/setup\_google\_colab.py
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.0.133, 151.101.64.133,
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101.0.133|:443... connect
HTTP request sent, awaiting response... 200 OK
Length: 1308 (1.3K) [text/plain]
Saving to: 'setup_google_colab.py'

setup_google_colab. 100%[=====>] 1.28K --.-KB/s in 0s

2019-06-12 09:06:23 (215 MB/s) - 'setup_google_colab.py' saved [1308/1308]
```

```
import tensorflow as tf
import keras
import numpy as np
import matplotlib.pyplot as plt
```

```
from keras.layers import Input, Dense, Lambda, InputLayer, concatenate
from keras.models import Model, Sequential
from keras import backend as K
from keras import metrics
from keras.datasets import mnist
from keras.utils import np_utils
from w5_grader import VAEGrader
```

```
[>] Using TensorFlow backend.
```

```
def vlb_binomial(x, x_decoded_mean, t_mean, t_log_var):
    """Returns the value of negative Variational Lower Bound

    The inputs are tf.Tensor
    x: (batch_size x number_of_pixels) matrix with one image per row with zeros and ones
    x_decoded_mean: (batch_size x number_of_pixels) mean of the distribution p(x | t), real numbers f
    t_mean: (batch_size x latent_dim) mean vector of the (normal) distribution q(t | x)
    t_log_var: (batch_size x latent_dim) logarithm of the variance vector of the (normal) distributio

    Returns:
    A tf.Tensor with one element (averaged across the batch), VLB
    """
    print(tf.shape(x))
    print(x.get_shape().as_list())
    print(x_decoded_mean.get_shape().as_list())
    print(t_mean.get_shape().as_list())
    print(t_log_var.get_shape().as_list())

    # Bernoulli: p^k * (1-p)^(1-k) for k in {0, 1}
    # k: ground truth
    # p: decoded value
    xent_loss = K.sum(x * K.log(x_decoded_mean + 1e-10) + (1 - x) * K.log(1 - x_decoded_mean + 1e-10), ax
    print(xent_loss.get_shape().as_list())

    kl_loss = 0.5 * K.sum((1 + t_log_var - K.square(t_mean) - K.exp(t_log_var)), axis=1)
    print(kl_loss.get_shape().as_list())

    # average across minibatch
    vae_loss = - K.mean(xent_loss + kl_loss)
    return vae_loss
```

```

# Start tf session so we can run code.
sess = tf.InteractiveSession()
# Connect keras to the created session.
K.set_session(sess)

batch_size = 100
original_dim = 784 # Number of pixels in MNIST images.
latent_dim = 10 #3 # d, dimensionality of the latent code t.
intermediate_dim = 256 # Size of the hidden layer.
epochs = 3

x = Input(batch_shape=(batch_size, original_dim))
x1 = Input(batch_shape=(1, original_dim))
def create_encoder(input_dim):
    # Encoder network.
    # We instantiate these layers separately so as to reuse them later
    encoder = Sequential(name='encoder')
    encoder.add(InputLayer([input_dim]))
    encoder.add(Dense(intermediate_dim, activation='relu'))
    encoder.add(Dense(2 * latent_dim))
    return encoder
encoder = create_encoder(original_dim)

get_t_mean = Lambda(lambda h: h[:, :latent_dim])
get_t_log_var = Lambda(lambda h: h[:, latent_dim:])
h = encoder(x)
t_mean = get_t_mean(h)
t_log_var = get_t_log_var(h)

h1 = encoder(x1)
t1_mean = get_t_mean(h1)
t1_log_var = get_t_log_var(h1)

# Sampling from the distribution
# q(t | x) = N(t_mean, exp(t_log_var))
# with reparametrization trick.
def sampling(args):
    """Returns sample from a distribution N(args[0], diag(args[1]))

    The sample should be computed with reparametrization trick.

    The inputs are tf.Tensor
    args[0]: (batch_size x latent_dim) mean of the desired distribution
    args[1]: (batch_size x latent_dim) logarithm of the variance vector of the desired distribution

    Returns:
    A tf.Tensor of size (batch_size x latent_dim), the samples.
    """
    t_mean, t_log_var = args
    epsilon = K.random_normal(shape=K.shape(t_mean), mean=0.0, stddev=1.0)
    return t_mean + K.exp(0.5 * t_log_var) * epsilon

t = Lambda(sampling)([t_mean, t_log_var])
t1 = Lambda(sampling)([t1_mean, t1_log_var])

def create_decoder(input_dim):
    # Decoder network
    # We instantiate these layers separately so as to reuse them later
    decoder = Sequential(name='decoder')
    decoder.add(InputLayer([input_dim]))
    decoder.add(Dense(intermediate_dim, activation='relu'))
    decoder.add(Dense(original_dim, activation='sigmoid'))
    return decoder
decoder = create_decoder(latent_dim)
x_decoded_mean = decoder(t)
x1_decoded_mean = decoder(t1)

loss = vlb_binomial(x, x_decoded_mean, t_mean, t_log_var)
vae = Model(x, x_decoded_mean)
# Keras will provide input (x) and output (x_decoded_mean) to the function that
# should construct loss, but since our function also depends on other
# things (e.g. t_means), it is easier to build the loss in advance and pass
# a function that always returns it.
vae.compile(optimizer=keras.optimizers.RMSprop(lr=0.001), loss=lambda x, y: loss)

```

```
↳ Tensor("Shape_6:0", shape=(2,), dtype=int32)
[100, 784]
[100, 784]
[100, 10]
[100, 10]
[100]
[100]
```

```
# train the VAE on MNIST digits
(x_train, y_train), (x_test, y_test) = mnist.load_data()
# One hot encoding.
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)


x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
```

```
epochs = 20
hist = vae.fit(x=x_train, y=x_train,
              shuffle=True,
              epochs=epochs,
              batch_size=batch_size,
              validation_data=(x_test, x_test),
              verbose=2)
```

```
↳
```

Epoch 1/20

Tensor("encoder_8/dense_18/BiasAdd:0", shape=(1, 20), dtype=float32)



```
fig = plt.figure(figsize=(10, 10))
for fid_idx, (data, title) in enumerate(
    zip([x_train, x_test], ['Train', 'Validation'])):
    n = 10 # figure with 10 x 2 digits
    digit_size = 28
    figure = np.zeros((digit_size * n, digit_size * 2))
    decoded = sess.run(x_decoded_mean, feed_dict={x: data[:batch_size, :]})
    for i in range(10):
        figure[i * digit_size: (i + 1) * digit_size,
            :digit_size] = data[i, :].reshape(digit_size, digit_size)
```

```

        figure[i * digit_size: (i + 1) * digit_size,
              digit_size:] = decoded[i, :].reshape(digit_size, digit_size)
    ax = fig.add_subplot(1, 2, fid_idx + 1)
    ax.imshow(figure, cmap='Greys_r')
    ax.set_title(title)
    ax.axis('off')
plt.show()

```



```

n_samples = 10 # To pass automatic grading please use at least 2 samples here.
# sampled_im_mean is a tf.Tensor of size 10 x 784 with 10 random
# images sampled from the vae model.
sample_t = K.random_normal(shape=(n_samples, latent_dim), mean=0.0, stddev=1.0)
sampled_im_mean = decoder(sample_t)

```

```

sampled_im_mean_np = sess.run(sampled_im_mean)
# Show the sampled images.
plt.figure()
for i in range(n_samples):
    ax = plt.subplot(n_samples // 5 + 1, 5, i + 1)
    plt.imshow(sampled_im_mean_np[i, :].reshape(28, 28), cmap='gray')
    ax.axis('off')
plt.show()

```

☞



```
# One-hot labels placeholder.
x = Input(batch_shape=(batch_size, original_dim))
label = Input(batch_shape=(batch_size, 10))

x1 = concatenate([x, label])

encoder = create_encoder(original_dim + 10)

h = encoder(x1)
cond_t_mean = get_t_mean(h)
cond_t_log_var = get_t_log_var(h)

print(x1.get_shape().as_list())
print(h.get_shape().as_list())
print(cond_t_mean.get_shape().as_list())
print(cond_t_log_var.get_shape().as_list())

t = Lambda(sampling)([cond_t_mean, cond_t_log_var])
t1 = concatenate([t, label])

decoder = create_decoder(latent_dim + 10)
cond_x_decoded_mean = decoder(t1)

#cond_t_mean = # Mean of the latent code (without label) for cvae model.
#cond_t_log_var = # Logarithm of the variance of the latent code (without label) for cvae model.
#cond_x_decoded_mean = # Final output of the cvae model.
```

```
[> [100, 794]
    [100, 20]
    [100, 10]
    [100, 10]
```

```
conditional_loss = vlb_binomial(x, cond_x_decoded_mean, cond_t_mean, cond_t_log_var)
cvae = Model([x, label], cond_x_decoded_mean)
cvae.compile(optimizer=keras.optimizers.RMSprop(lr=0.001), loss=lambda x, y: conditional_loss)
```

```
[> Tensor("Shape_7:0", shape=(2,), dtype=int32)
    [100, 784]
    [100, 784]
    [100, 10]
    [100, 10]
    [100]
    [100]
```

```
hist = cvae.fit(x=[x_train, y_train],
                y=x_train,
                shuffle=True,
                epochs=epochs,
                batch_size=batch_size,
                validation_data=([x_test, y_test], x_test),
                verbose=2)
```

```
[>
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

- 2s - loss: 155.4058 - val_loss: 129.0842

Epoch 2/20

- 2s - loss: 124.1803 - val_loss: 120.4274

Epoch 3/20

- 2s - loss: 118.4696 - val_loss: 116.5869

Epoch 4/20

- 2s - loss: 115.4052 - val_loss: 113.4894

Epoch 5/20

- 2s - loss: 113.3610 - val_loss: 111.9370

Epoch 6/20

- 2s - loss: 111.9178 - val_loss: 110.4100

Epoch 7/20

- 2s - loss: 110.7301 - val_loss: 109.6823

Epoch 8/20

- 2s - loss: 109.8311 - val_loss: 108.2383

Epoch 9/20

- 2s - loss: 109.0987 - val_loss: 107.7660

Epoch 10/20

- 2s - loss: 108.4247 - val_loss: 108.0507

Epoch 11/20

- 2s - loss: 107.8561 - val_loss: 106.8372

Epoch 12/20

- 2s - loss: 107.3355 - val_loss: 106.8308

Epoch 13/20

- 2s - loss: 106.8826 - val_loss: 105.6550

```
fig = plt.figure(figsize=(10, 10))
for fid_idx, (x_data, y_data, title) in enumerate(
    zip([x_train, x_test], [y_train, y_test], ['Train', 'Validation'])):
    n = 10 # figure with 10 x 2 digits
    digit_size = 28
    figure = np.zeros((digit_size * n, digit_size * 2))
    decoded = sess.run(cond_x_decoded_mean,
                       feed_dict={x: x_data[:batch_size, :],
                                   label: y_data[:batch_size, :]})

    for i in range(10):
        figure[i * digit_size: (i + 1) * digit_size,
              :digit_size] = x_data[i, :].reshape(digit_size, digit_size)
        figure[i * digit_size: (i + 1) * digit_size,
              digit_size:] = decoded[i, :].reshape(digit_size, digit_size)
    ax = fig.add_subplot(1, 2, fid_idx + 1)
    ax.imshow(figure, cmap='Greys_r')
    ax.set_title(title)
    ax.axis('off')
plt.show()
```





```
# Prepare one hot labels of form
# 0 0 0 0 0 1 1 1 1 1 2 2 2 2 2 ...
# to sample five zeros, five ones, etc
curr_labels = np.eye(10)
curr_labels = np.repeat(curr_labels, 5, axis=0) # Its shape is 50 x 10.
# cond_sampled_im_mean is a tf.Tensor of size 50 x 784 with 5 random zeros,
# then 5 random ones, etc sampled from the cvae model.
sample_t = K.random_normal(shape=(50, latent_dim), mean=0.0, stddev=0.5)
sample_tl = concatenate([sample_t, tf.convert_to_tensor(curr_labels, dtype=tf.float32)])
cond_sampled_im_mean = decoder(sample_tl)
```

```
cond_sampled_im_mean_np = sess.run(cond_sampled_im_mean)
# Show the sampled images.
plt.figure(figsize=(10, 10))
global_idx = 0
for digit in range(10):
    for _ in range(5):
        ax = plt.subplot(10, 5, global_idx + 1)
        plt.imshow(cond_sampled_im_mean_np[global_idx, :].reshape(28, 28), cmap='gray')
        ax.axis('off')
        global_idx += 1
plt.show()
```

```
from keras.datasets import cifar10
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

↳ Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 6s 0us/step

```
plt.imshow(x_train[7, :])
plt.show()
```

